

Distributed Deep Learning on GPU-based Clusters

Prajwal Singhanian, Lannie Dalton Hough, Abhinav Bhatele
Department of Computer Science

Data parallelism

- Work on different parts of the data in parallel on different GPUs
- Example: PyTorch's DDP, FSDP, and DeepSpeed's ZeRO

Work distribution in data parallelism

- Each worker has a full copy of the entire NN and processes different shards of batch
- All reduce operation to synchronize gradients



DDP Strategy in Lightning

```
from lightning.fabric.strategies import DDPStrategy

pl_strategy = DDPStrategy()

fabric = Fabric(
    strategy=pl_strategy,
    . . .)
```

Running the code

- Code location - train.py

```
CONFIG_FILE=configs/ddp.json sbatch train_multi.sh
```

FSDP

- Using DDP is limited to smaller model sizes
- FSDP implements memory optimizations to fit larger models on a GPU
 - Shard parameters of each Decoder block across GPUs

FSDPStrategy in Lightning

```
from lightning.fabric.strategies import
FSDPStrategy


pl_strategy = FSDPStrategy(
    auto_wrap_policy={Block}
)

fabric = Fabric(strategy=pl_strategy, . . .)
```

Specify what
modules to
shard

Running the code

```
CONFIG_FILE=configs/fsdp.json sbatch train_multi.sh
```



Prajwal Singhania, Lannie Dalton Hough, Abhinav Bhatele

Department of Computer Science

prajwal@umd.edu, bhatele@cs.umd.edu, ldhough@umd.edu

